

Highlights

Lead-biased Grounding Score: injecting discourse-structural knowledge into reference-free, cost-constrained assessment of summary quality

Mikolaj Semeniuk, Agnieszka Jastrzebska

- LGS scores summary quality with no human reference and no LLM judge call
- Discourse-position knowledge enters as a one-parameter prior over source sentences
- Only 5 of 14 metrics are Pareto-optimal in cost-quality; LGS is one of them
- LGS beats BERTScore in mean correlation at one tenth of the runtime
- The knowledge prior must be recalibrated for each sentence-embedding backbone